

## **Homework 1:** ***Supervised Learning***

### **Datasets and why are they important?**

#### **Nursery**

The first dataset I have chosen comes from 1980's Ljubljana, Slovenia. During that period, there used to be an excessive number of applications to nursery schools but not enough school capacity to admit all of the applicants. Therefore, the government discretized some certain aspects of the applicants by assigning points and discrete values to them. In the end, there were a total of three main categories that the applications were evaluated: parents' occupations, family structure and financial standing, and social and health picture of the family. In this dataset, these three subcategories are described with 8 different features (attributes) and these attributes have around 3-5 options each. The target value (label) is also categorized in five different categories, ranging from '*not recommended*' to '*special priority*'. In this dataset, there are 12960 instances, meaning 12960 children who are classified. I separated my training and test data as 80% and 20% respectively of the total number of instances.

The reason I have chosen this dataset is because it comes from real life and the end result (labels) was something that had a great impact on thousands of families and their children. In reality, government formulized a way to classify the children and in this project, I will be trying to capture their classification methods and models. Throughout this paper, I will be referring this dataset as the 'nursery dataset.'

#### **Letter Recognition (Letters)**

This dataset contains 20,000 instances of digitally typed and distorted in different ways letters from 20 different fonts of the English alphabet, meaning the result (label) has 26 different available options. Each letter appears between 734-813 times in the dataset, suggesting a somewhat balanced distribution. Each instance is defined by 16 different features including height, width of the char box, the number of pixels available and other features that describe edges of the letters. Each attribute is discredited from a continuous scale to an integer range between 1-15.

The reason I have chosen this dataset was to be able to compare and contrast with the first dataset. Since this dataset contains attributes that are derived from a continuous range, intuitively, the classification process seems to be different compared to the first one, a dataset with very discreet and polarized features. Another reason I have chosen this dataset was to be able to work with one of the most popular concepts in the industry: OCR (optical character recognition). I wanted to see how different machine learning algorithms would be able to capture the basics of this method and be successful when it comes to recognizing shapes and classifying them as the correct letters. Throughout this paper, I will be referring this dataset as the 'letter dataset.'

## Project Overview

This homework will have analyses of 5 different supervised learning algorithms. I will be exploring how different variables and kernels affect the results of these algorithms and how do they perform when ran on the two datasets I have. I separated each dataset as 80% training data, and 20% testing data to be able to test the trained classifiers on the test data. In the analyses part, there will be learning curve graphs and varying characteristic variables versus the accuracy graphs. Most of the classifier models are cross-validated 7 times and this will be later explored.

## Decision Trees

Decision trees are one of the most popular supervised learning algorithms that build as classifier tree using the simple decisions inferred from the training data. In order to experiment with decision trees I tried out algorithms both in Weka and python's sklearn library.

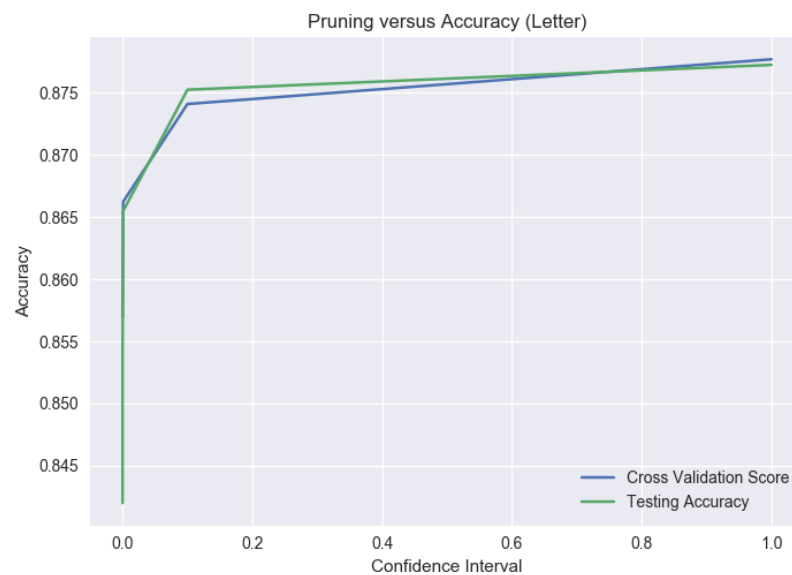
### Nursery



**Figure 1:** Learning curve for decision tree on the nursery dataset, no pruning is used in the decision tree/

The *figure 1* above shows the learning curve for the decision tree algorithm applied on the nursery dataset. As it can be seen, as the training set size reaches to 100%, the accuracy of both the testing and cross validation sets are increasing. This shows that decision trees perform better as there is more and more data points to learn from. Throughout the rise of the both curves, we notice that cross-validation score and the testing scores are quite similar and react the same way to the increase in the training set size, this shows that data (labels in this situation) is evenly distributed. Since there's no pruning in this decision tree, it performs almost 100% accurate when the entire training set is used. The effects of pruning will be discussed in the next dataset's example.

## Letter Recognition

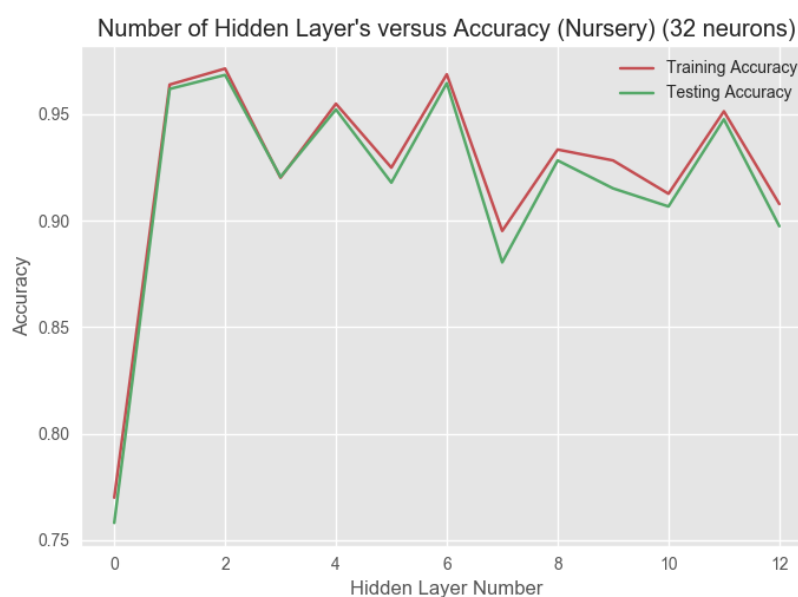


**Figure 2:** Model complexity graph, showing effects of pruning on decision trees on the letter dataset

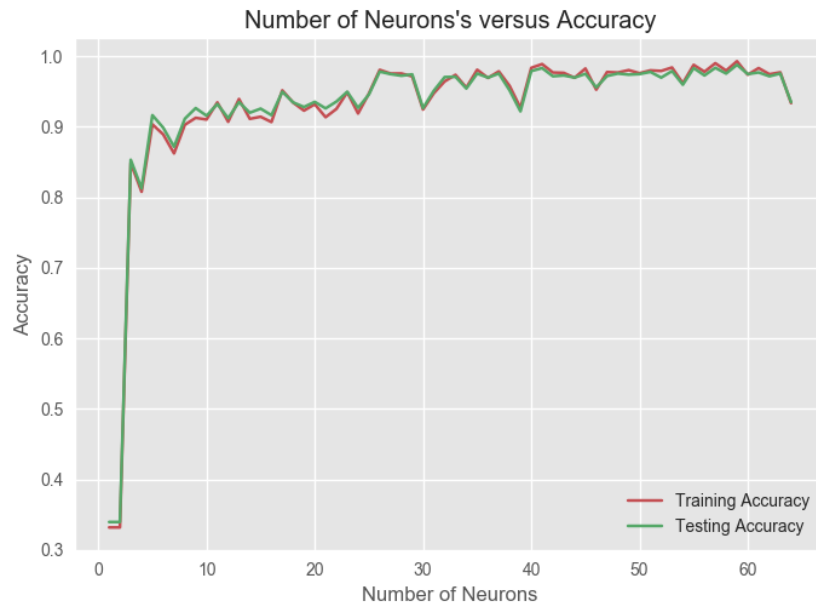
The *figure 2* above shows how decreasing the pruning (increasing the confidence level) increases the accuracy on decision trees. As the confidence interval increases, the decision tree is less and less pruned, meaning that it's allowed to be taller and deeper. However, the reason that we prune is to lose the parts of the tree that has little effects on the classification process, therefore getting rid of noise and preventing any overfitting that might happen because of that. An increasing accuracy with decreasing pruning levels mean that letter dataset might not have too much noise, in other words, the variance of the dataset might be low.

## Neural Nets

For the neural nets, I ran experiments on the number of hidden layers and the number of neurons in each hidden layer, as well as changing the training set size to see if neurons need a lot of data and examples as it seems intuitively.



**Figure 3:** Model complexity graph for neural nets on nursery dataset, exploring hidden layer number



**Figure 4:** The model complexity graph for neural nets on nursery dataset, exploring neuron number.

Figures 3 and 4 show how changing the hidden layer number and neuron number on each hidden layer affects the learning. From the *figure 3*, it can be inferred that 2 is the optimal hidden layer number for the dataset even though 4 layers produce the same result, because if accuracies are the same, the simplest one should always be preferred. In *figure 4*, it can be seen that as the neuron number on each layer increases, the accuracies increase. Around 60 neurons, the classifier predicts almost all of the test data and the training data right. I have not applied cross validation in my neural net algorithm since building a model was taking already too long. Cross validating in the training set would building the classifier model each time from scratch, making this algorithm implausible for this homework's sake.

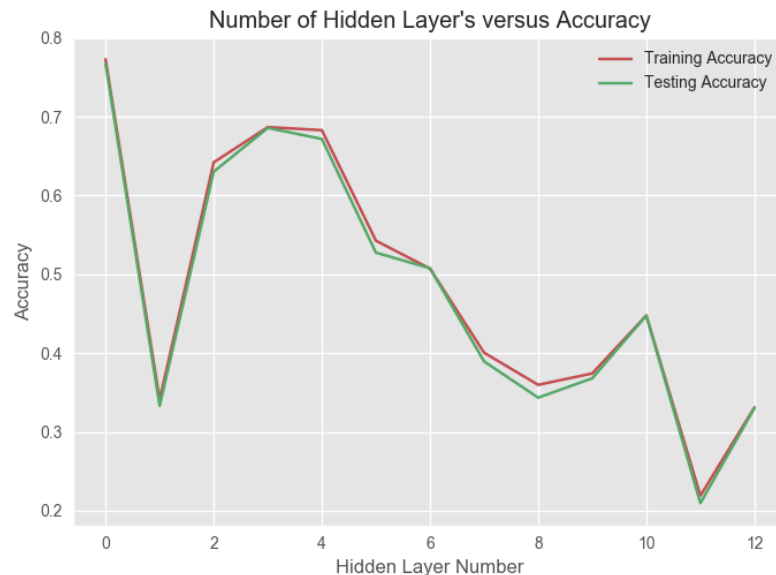


**Figure 5:** Learning curve for neural nets on nursery dataset (2 hidden layers, 60 neurons).

The *figure 5* shows the learning curve for neural nets on the nursery dataset. This graph summarizes how neural networks very well. It can be seen that on lower training set sizes, there's a significant difference between training and test accuracies and training accuracy surpasses the testing accuracy. This kind of a start suggests that the model still underfits the

training data, creating a classifier good enough for seen data but not enough for unseen data. As the number of training data increases, two lines start to merge and converge together suggesting that the results they obtain become closer to each other.

## Letter Recognition



**Figure 6:** Model complexity graph for neural nets on the letter dataset, exploring hidden layer number

Trying different numbers of hidden values, I came to the conclusion that 0 hidden layers work the best for the letter dataset when it comes to neural networks. As the *figure 6*, suggests, adding more and more hidden layers decrease the accuracy from around 80% to almost 20% when the hidden layer value becomes 10. I know that even adding 2 hidden layers would be too much for such a small dataset but since the computation did not take as much, I explored until 12. The *figure 7* below shows the learning curve for this algorithm and as it can be seen, the plot shares a lot of similarities with the previous learning curve for the neural net learning curve on the nursery dataset. On smaller training sets underfitting is an issue and as the data number to train increases, the model starts performing better and training and testing accuracies come closer.

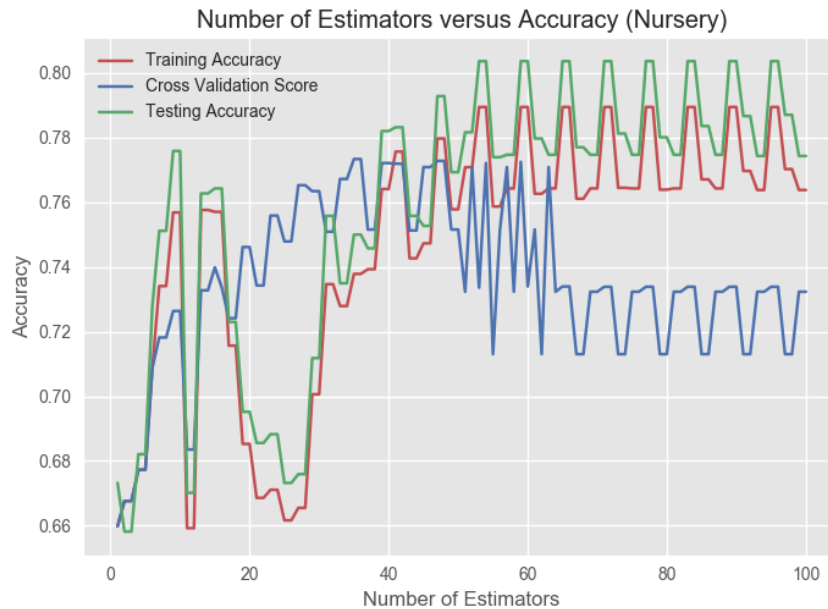


**Figure 7:** Learning curve for neural nets on the letter dataset

## Boosting

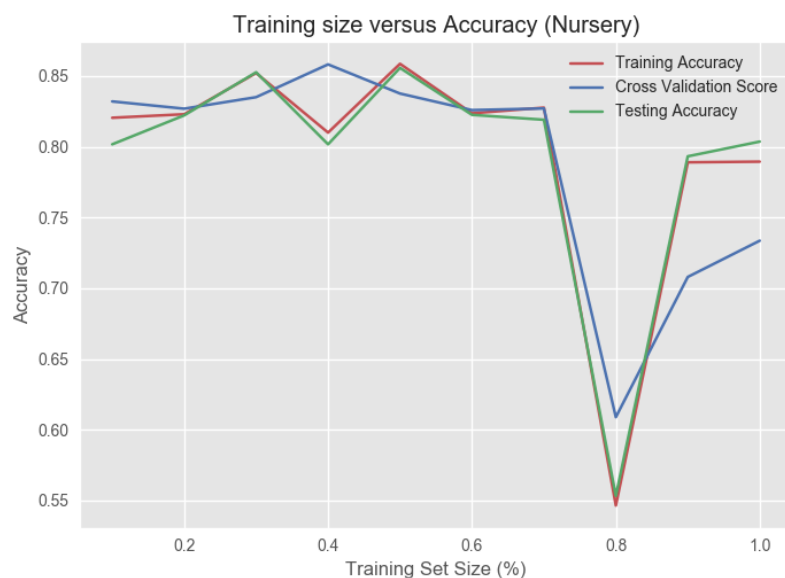
For boosting, I used python's Adaboost algorithm. Adaboost is an algorithm that performs really well on binary classifications or with classification that doesn't have too many labels. The effects of that on both of my algorithms will be discussed in the section. I also experimented with changing the number of weak classifiers (estimators) to start with.

### Nursery



**Figure 8:** Changing the number of estimators in boosting applied on the nursery dataset.

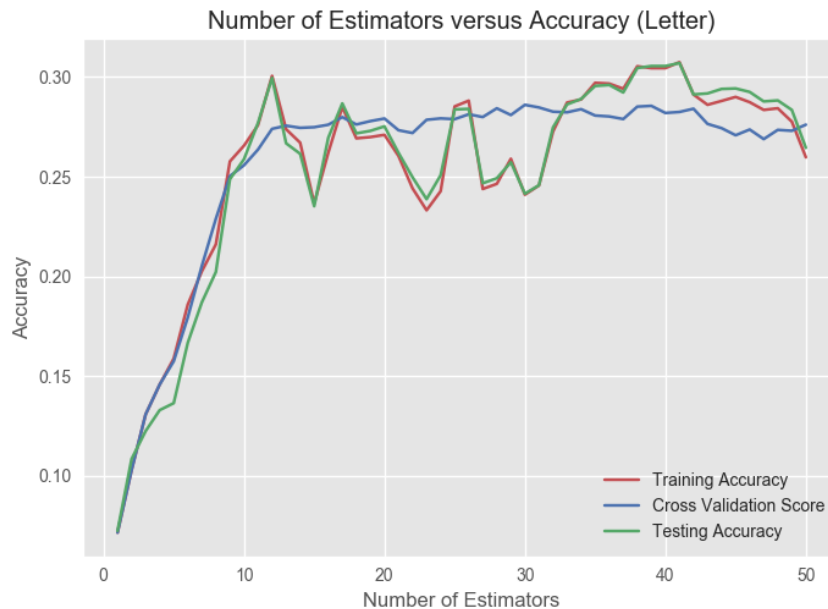
The graph above in *figure 8* shows the effects of changing the number of estimators in boosting. Surely, it looks much different than any other graphs on this paper. The cross validation score shows here that around 35 estimators where the accuracy is peaking for the number of estimators, and after that it drops, causing the model to overfit due to having a high number of trees (high number of estimators) and dealing with a lot of noise. Also, almost a wave-like curve can be observed in all three curves, suggesting that the modulo of the number of estimators might have an effect on the learning itself.



**Figure 9:** Learning curve for boosting on the nursery dataset

The learning curve of boosting on the nursery dataset suggests that the training set size doesn't have a direct effect on the accuracy. In fact, accuracy (if the discrepancy at 0.8 point on x is not considered) is slowly decreasing. This shows that boosting doesn't depend on the training data that much.

## Letter Recognition



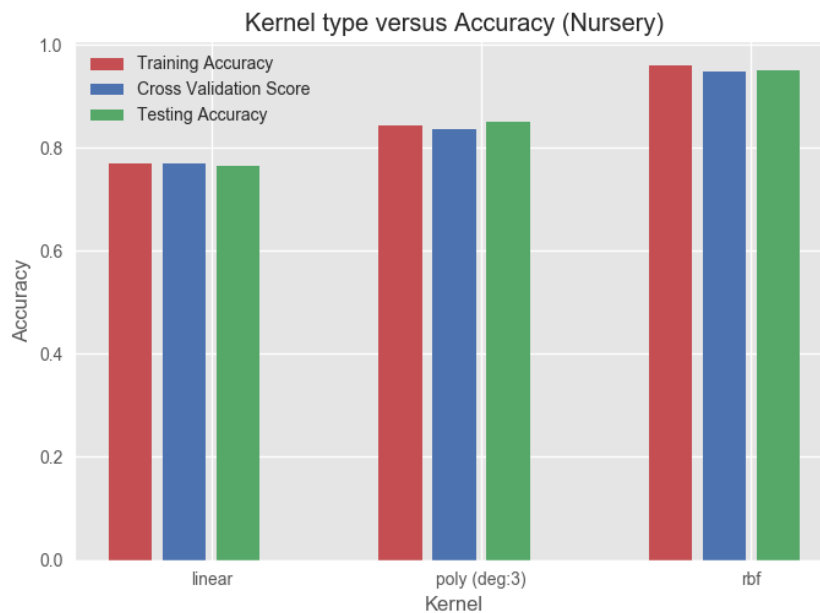
**Figure 10:** Model complexity graph for the number of estimators in boosting (letter dataset)

The *figure 10* above shows that the optimal number of estimators for the letter dataset for boosting is around 12. However, it can be stated that at the peak, the accuracy is only around 30%, suggesting that boosting performs really poorly on this dataset. The reason for that is the nature of adaboost since it is designed for either binary classifications or low number of labels. Since the letter dataset has 26 labels, it is understandable to have such low level of accuracy. One way to prevent this was to choose one label and mark it and mark the rest of the labels to be another label, reducing the dataset to a binary classification and then repeat this for each label to find the boosting performance. The learning curve for boosting shows around 30% accuracy for most of the training set sizes. It doesn't point out anything interesting about this algorithm and therefore, it's not shown in this paper.

## Support Vector Machine

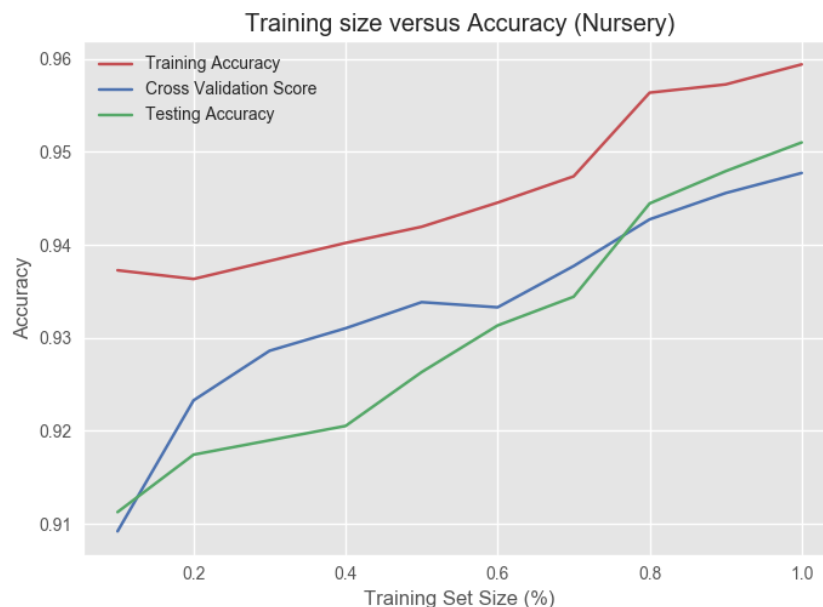
Support vector machine algorithms use different kernels to separate the instances into clusters. In my experiments with SVM's, I explored different kernels and how each one performs on my datasets, as well as experimenting how the training set size affects the accuracy.

## Nursery



**Figure 11:** The model complexity chart for SVM on the nursery dataset that explores different kernels.

The bar chart given in *figure 11* shows how different kernels of SVM work on the nursery dataset. As it can be seen, as the complexity of the kernel is increasing, the accuracy also increases for this dataset. While experimenting, the linear kernel took around half the time of the radial basis function kernel, suggesting that the complexity has a direct effect on the running time. The RBF kernel performs with 96.5% accuracy on this dataset while the linear kernel performs only with 78% percent.

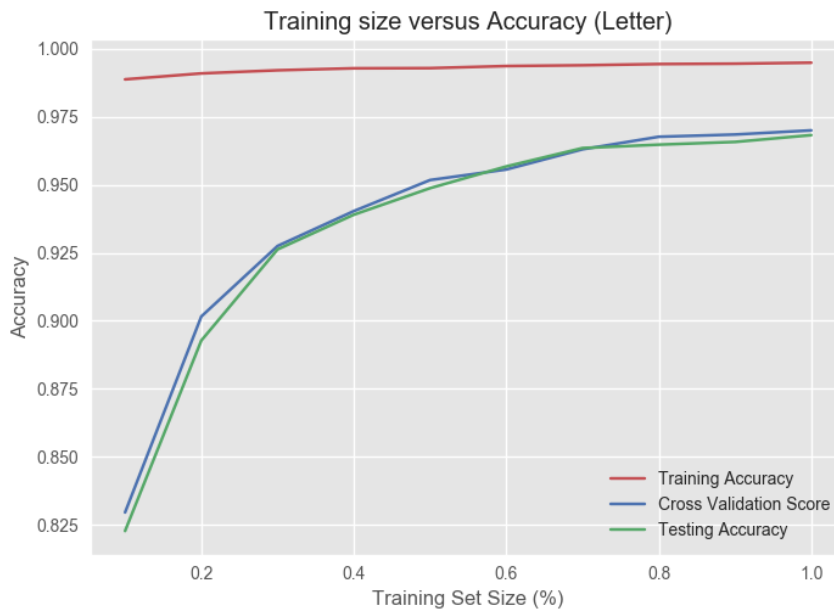


**Figure 12:** Learning curve for SVM on the nursery dataset, using RBF kernel

The *figure 12* shows the learning curve for SVM. The red line (training accuracy) is consistently 2-3% higher than the green line, which is the testing accuracy. This suggests that overfitting is an issue here but not very significantly. Also, even though the 10% of the entire training set given, the algorithm starts to perform with a little more than 91% accuracy, revealing that unlike most of the algorithms, SVM does not highly depend on the training data. That's because the SVM creates an estimator to classify and usually uses only relevant data to separate clusters.



## Letter Recognition



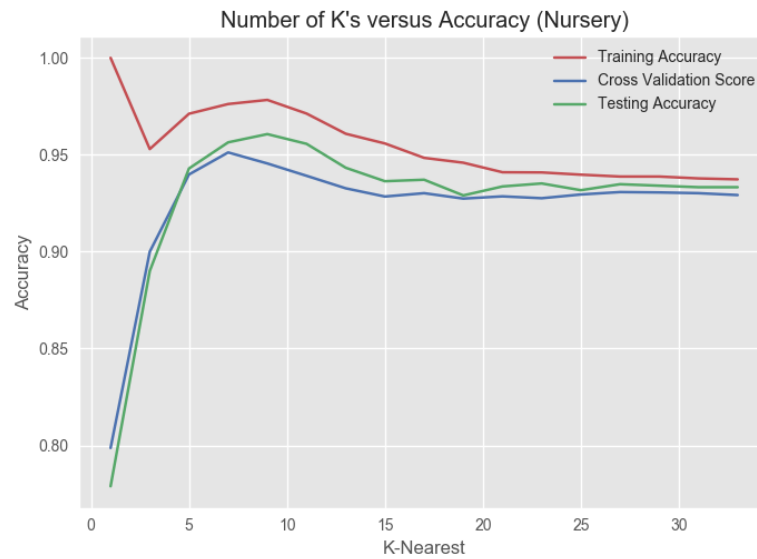
**Figure 13:** Learning curve for SVM on the letter dataset, using the RBF kernel

After testing different kernels on the letter dataset, checking the cross validation score and seeing how close it is to the testing score on RBF kernel as well as its high accuracy value, I decided that RBF was the optimal kernel for this dataset as well. The *figure 13* shows the learning curve for SVM. The most striking feature of this graph is the very high and consistent training accuracy curve. Compared to that, cross validation and testing accuracies start with a lower accuracy and gets closer to the training curve, however, cannot catch it. This suggests that the fitted SVM model overfits the testing data. Since the increasing training set size suggests that the more the data, the less overfitting, we can come to the conclusion that the complexity of the kernel might be overcomplicating the function. Another kernel with lesser complexity could have performed better on this dataset.

## K-Nearest Neighbors

K-nearest neighbors is a classification algorithm that tries to classify the instance by finding the closest k-number of neighbors and identifying that the instance belongs to that cluster. The first step I have taken was to find the optimal k value for each dataset and then I plotted learning curves to see if changing the training set size has a significant effect on the accuracy.

## Nursery



**Figure 14:** The model complexity graph for KNN algorithm on nursery dataset for k values.

Running the KNN algorithm on the nursery data set using a range of odd k values between 1 and 35, I obtained the model complexity graph given in *figure 14*. Looking at the graph, it can be seen that on lower k values, the trained model performs really well on the trained data, scoring close to 100% and over 95%. On the other hand, at these k values, the testing accuracy is low, suggesting the model performs much better on the seen data compared to unseen data. This is a typical example of overfitting since the model tailors its classifier very specifically for the trained data. The plot also reveals that the optimum k value for the nursery dataset is 9, and after this point, overfitting rate decreases and all three accuracy values start to converge but decrease, suggesting that increasing the k value all the way doesn't increase the accuracy. In the end, looking at a limited amount of neighbors only make sense and as k increases, the algorithm starts to merge different clusters into each other, decreasing the accuracy and causing false positives.

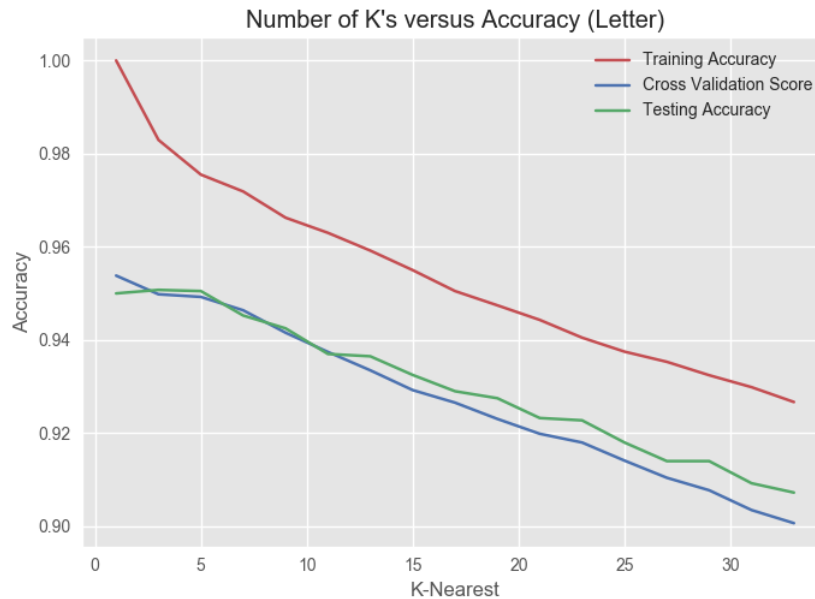


**Figure 15:** Learning curve of KNN on the nursery dataset

Looking at the *figure 15*, it can be observed that as the training set size increases, the accuracy of the model increases as well. However, it is possible to observe that the training accuracy is always higher than the testing accuracy, showing that the algorithm performs better on the

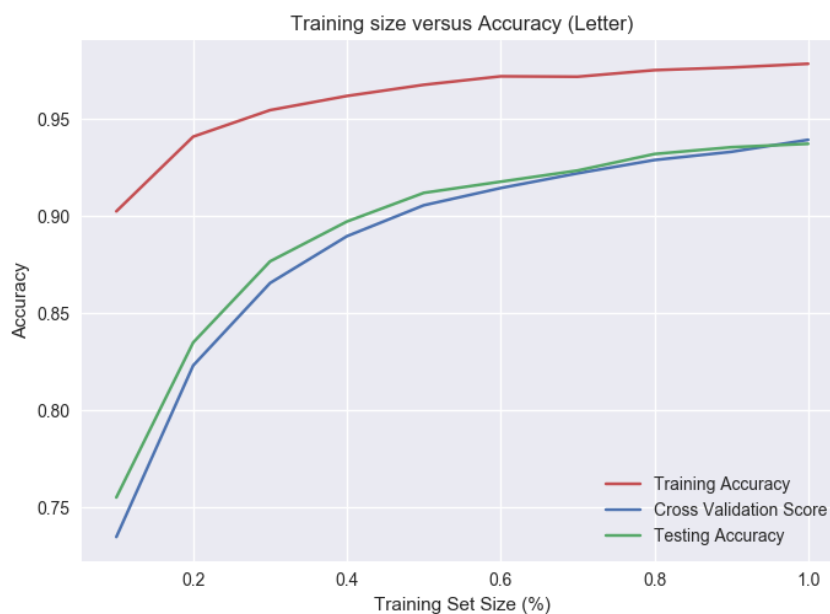
seen data. A viable reason for this occasion is the KNN algorithm uses all the given training data rather than using the data to make an estimator, therefore, the entire training data is stored to run this algorithm, which leads to overfitting, creating complex separators that consider noise as real data. Also, having very similar testing and cross validation curves show that there is not much bias in the testing set. In the final analysis, we see that, given the entire training set, the model performs with around 93.5% accuracy, suggesting that with 9-nearest neighbors, the labels can be classified and clustered with a significant accuracy.

## Letter Recognition



**Figure 16:** The model complexity graph for KNN algorithm on the letter dataset for k's

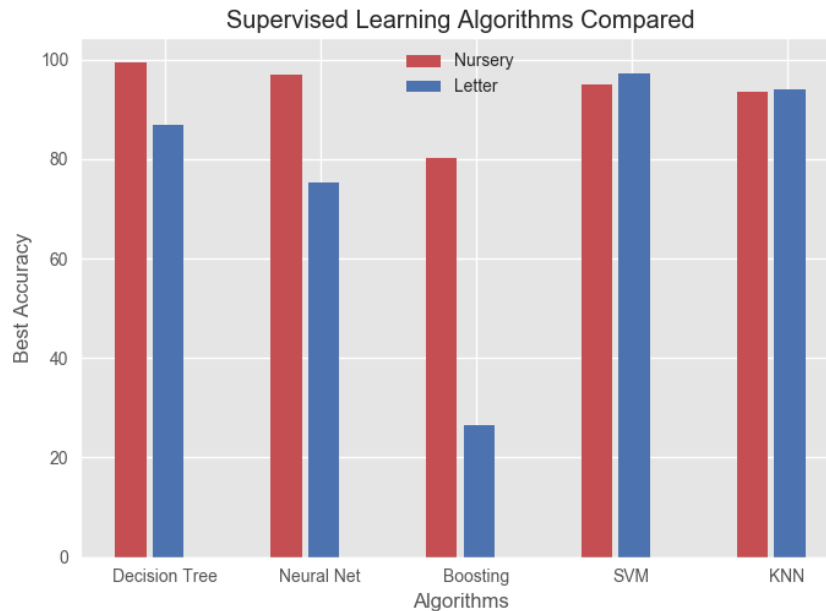
The *figure 16* shows the accuracy for KNN algorithm on the letter dataset, iterating over different k values ranging from 1 to 35. Unlike the previous dataset, the k value starts decreasing immediately, suggesting the optimum k value is around 3.



**Figure 17:** Learning curve for KNN algorithm on the letter dataset.

The learning curve in the *figure 17* suggests that KNN has an increasing accuracy with the increasing training sample set size. This makes sense since KNN stores all the input instances, meaning the more possible neighbors, the more accurate the model gets. Other than that, overfitting is an issue like the previous dataset, since training accuracy exceeds the testing and cross validation data in a significant way.

## Conclusion



**Figure 18:** Comparison for all 5 supervised algorithms, for each of the datasets. This bar chart is created considering only the best accuracies with the highest training set sizes.

In the final analysis, *figure 18* shows how each algorithm performed on each of the datasets I have chosen. It can be seen that the nursery dataset either has a very similar or higher accuracy values compared to the letter dataset. This can be explained with the number of labels each dataset has. The letter dataset was much harder to classify with 26 labels when the nursery dataset had only around 5. It can also be stated that KNN and SVM algorithms performed around the same for each dataset with relatively higher accuracy percentages, which are around 95 percent. The best performing algorithm was decision trees for the nursery dataset and this can be explained with the low-variance nature of the dataset itself. Without any pruning, the algorithm performed with better than 99% accuracy. The most unsuccessful algorithm was boosting since the specific algorithm I used was better for binary classifications, and therefore I should have chose a better boosting algorithm to capture better classification models of my datasets.